

# mmap – Memory-mapped file support

Memory-mapped file objects behave like both bytearray and like file objects. You can use mmap objects in most places where bytearray are expected; for example, you can use the re module to search through a memory-mapped file. You can also change a single byte by doing `obj[index] = 97`, or change a subsequence by assigning to a slice: `obj[i1:i2] = b'...'`. You can also read and write data starting at the current file position, and `seek()` through the file to different positions.

A memory-mapped file is created by the mmap constructor, which is different on Unix and on Windows. In either case you must provide a file descriptor for a file opened for update. If you wish to map an existing Python file object, use its `fileno()` method to obtain the correct value for the `fileno` parameter. Otherwise, you can open the file using the `os.open()` function, which returns a file descriptor directly (the file still needs to be closed when done).

To map anonymous memory, -1 should be passed as the `fileno` along with the length.

```
class mmap.mmap(fileno, length, tagname=None, access=ACCESS_DEFAULT[, offset])
```

(Windows version) Maps `length` bytes from the file specified by the file handle `fileno`, and creates a mmap object. If `length` is larger than the current size of the file, the file is extended to contain `length` bytes. If `length` is 0, the maximum length of the map is the current size of the file, except that if the file is empty Windows raises an exception (you cannot create an empty mapping on Windows).

`tagname`, if specified and not `None`, is a string giving a tag name for the mapping. Windows allows you to have many different mappings against the same file. If you specify the name of an existing tag, that tag is opened, otherwise a new tag of this name is created. If this parameter is omitted or `None`, the mapping is created without a name. Avoiding the use of the tag parameter will assist in keeping your code portable between Unix and Windows.

`offset` may be specified as a non-negative integer offset. mmap references will be relative to the offset from the beginning of the file. `offset` defaults to 0. `offset` must be a multiple of the `ALLOCATIONGRANULARITY`.

Raises an auditing event `mmap.__new__` with arguments `fileno`, `length`, `access`, `offset`.

```
class mmap.mmap(fileno, length, flags=MAP_SHARED, prot=PROT_WRITE|PROT_READ,
access=ACCESS_DEFAULT[, offset])
```

(Unix version) Maps `length` bytes from the file specified by the file descriptor `fileno`, and returns a mmap object. If `length` is 0, the maximum length of the map will be the current size of the file when mmap is called.

flags specifies the nature of the mapping. MAP\_PRIVATE creates a private copy-on-write mapping, so changes to the contents of the mmap object will be private to this process, and MAP\_SHARED creates a mapping that's shared with all other processes mapping the same areas of the file. The default value is MAP\_SHARED.

prot, if specified, gives the desired memory protection; the two most useful values are PROT\_READ and PROT\_WRITE, to specify that the pages may be read or written. prot defaults to PROT\_READ | PROT\_WRITE.

access may be specified in lieu of flags and prot as an optional keyword parameter. It is an error to specify both flags, prot and access. See the description of access above for information on how to use this parameter.

offset may be specified as a non-negative integer offset. mmap references will be relative to the offset from the beginning of the file. offset defaults to 0. offset must be a multiple of ALLOCATIONGRANULARITY which is equal to PAGE\_SIZE on Unix systems.

To ensure validity of the created memory mapping the file specified by the descriptor fileno is internally automatically synchronized with physical backing store on Mac OS X and OpenVMS.

This example shows a simple way of using mmap:

**import mmap** # write a simple example file

**with open("hello.txt", "wb") as f:**

```
    f.write(b"Hello Python!\n") with open("hello.txt", "r+b") as f: # memory-map the file,
    size 0 means whole file mm = mmap.mmap(f.fileno(), 0) # read content via standard
    file methods print(mm.readline())
```

```
    # prints b"Hello Python!\n"
```

```
    # read content via slice notation
```

```
    print(mm[:5])
```

```
    # prints b"Hello"
```

```
    # update content using slice notation;
```

```
    # note that new content must have same size
```

```
    mm[6:] = b" world!\n"
```

```
    # ... and read again using standard file methods
```

```
    mm.seek(0)
```

```
    print(mm.readline())
```

```
# prints b"Hello world!\n"
```

```
# close the map
```

```
mm.close()
```

mmap can also be used as a context manager in a with statement:

```
import mmap
```

```
with mmap.mmap(-1, 13) as mm:
```

```
    mm.write(b"Hello world!")
```

The next example demonstrates how to create an anonymous map and exchange data between the parent and child processes:

```
import mmap
```

```
import osmm = mmap.mmap(-1, 13)
```

```
mm.write(b"Hello world!")pid = os.fork()if pid == 0:
```

```
    # In a child process
```

```
    mm.seek(0)
```

```
    print(mm.readline())
```

```
    mm.close()
```

Raises an auditing event `mmap.__new__` with arguments `fileno`, `length`, `access`, `offset`.

Memory-mapped file objects support the following methods:

`close()`

Closes the mmap. Subsequent calls to other methods of the object will result in a `ValueError` exception being raised. This will not close the open file.

`closed`

True if the file is closed.

`find(sub[, start[, end]])`

Returns the lowest index in the object where the subsequence `sub` is found, such that `sub` is contained in the range `[start, end]`. Optional arguments `start` and `end` are interpreted as in slice notation. Returns `-1` on failure.

`flush([offset[, size]])`

Flushes changes made to the in-memory copy of a file back to disk. Without use of this call there is no guarantee that changes are written back before the object is destroyed. If offset and size are specified, only changes to the given range of bytes will be flushed to disk; otherwise, the whole extent of the mapping is flushed. offset must be a multiple of the PAGESIZE or ALLOCATIONGRANULARITY.

`madvise(option[, start[, length]])`

Send advice option to the kernel about the memory region beginning at start and extending length bytes. option must be one of the MADV\_\* constants available on the system. If start and length are omitted, the entire mapping is spanned. On some systems (including Linux), start must be a multiple of the PAGESIZE.

`move(dest, src, count)`

Copy the count bytes starting at offset src to the destination index dest. If the mmap was created with ACCESS\_READ, then calls to move will raise a TypeError exception.

`read([n])`

Return a bytes containing up to n bytes starting from the current file position. If the argument is omitted, None or negative, return all bytes from the current file position to the end of the mapping. The file position is updated to point after the bytes that were returned.

`read_byte()`

Returns a byte at the current file position as an integer, and advances the file position by 1.

`readline()`

Returns a single line, starting at the current file position and up to the next newline. The file position is updated to point after the bytes that were returned.

`resize(newsize)`

Resizes the map and the underlying file, if any. If the mmap was created with ACCESS\_READ or ACCESS\_COPY, resizing the map will raise a TypeError exception.

`rfind(sub[, start[, end]])`

Returns the highest index in the object where the subsequence sub is found, such that sub is contained in the range [start, end]. Optional arguments start and end are interpreted as in slice notation. Returns -1 on failure.

`seek(pos[, whence])`

Set the file's current position. whence argument is optional and defaults to os.SEEK\_SET or 0 (absolute file positioning); other values are os.SEEK\_CUR or 1 (seek relative to the current position) and os.SEEK\_END or 2 (seek relative to the file's end).

`size()`

Return the length of the file, which can be larger than the size of the memory-mapped area.

`tell()`

Returns the current position of the file pointer.

`write(bytes)`

Write the bytes in bytes into memory at the current position of the file pointer and return the number of bytes written (never less than len(bytes), since if the write fails, a ValueError will be raised). The file position is updated to point after the bytes that were written. If the mmap was created with ACCESS\_READ, then writing to it will raise a TypeError exception.

`write_byte(byte)`

Write the integer byte into memory at the current position of the file pointer; the file position is advanced by 1. If the mmap was created with ACCESS\_READ, then writing to it will raise a TypeError exception.

`MADV_* Constants`

`mmap.MADV_NORMAL`

`mmap.MADV_RANDOM`

`mmap.MADV_SEQUENTIAL`

`mmap.MADV_WILLNEED`

`mmap.MADV_DONTNEED`

`mmap.MADV_REMOVE`

`mmap.MADV_DONTFORK`

`mmap.MADV_DOFORK`

`mmap.MADV_HWPOISON`

`mmap.MADV_MERGEABLE`

`mmap.MADV_UNMERGEABLE`

`mmap.MADV_SOFT_OFFLINE`

mmap.MADV\_HUGEPAGE  
mmap.MADV\_NOHUGEPAGE  
mmap.MADV\_DONTDUMP  
mmap.MADV\_DODUMP  
mmap.MADV\_FREE  
mmap.MADV\_NOSYNC  
mmap.MADV\_AUTOSYNC  
mmap.MADV\_NOCORE  
mmap.MADV\_CORE  
mmap.MADV\_PROTECT

These options can be passed to `mmap.madvise()`. Not every option will be present on every system.